

# Understanding Cosine Similarity and Its Use in Vector Databases

## 1. Idea

Each object (text, image, recipe, etc.) can be represented as a list of numbers called a **vector**. Cosine similarity measures how much two vectors point in the same direction in space. If they point the same way → similarity = 1. If they are perpendicular → 0. If they go opposite → -1.

## 2. The formula

$$\text{Cosine Similarity}(A, B) = (A \cdot B) / (\|A\| \times \|B\|)$$

**Dot product (A · B):** multiply each component and add them up.

**Norm ( $\|A\|$ ):** the length of the vector, calculated as the square root of the sum of squares.

**Division:** divides the dot product by the product of the lengths to normalize.

### Example:

$$A = [1, 2, 3], B = [2, 4, 6]$$

$$A \cdot B = 1 \cdot 2 + 2 \cdot 4 + 3 \cdot 6 = 28$$

$$\|A\| = \sqrt{(1^2 + 2^2 + 3^2)} = \sqrt{14} \approx 3.74$$

$$\|B\| = \sqrt{(2^2 + 4^2 + 6^2)} = \sqrt{56} \approx 7.48$$

$$\text{Cosine similarity} = 28 / (3.74 \times 7.48) = 1 \rightarrow \text{perfectly aligned.}$$

### Intuitive meaning:

1 → Vectors go in the same direction (very similar)

0 → Unrelated vectors

-1 → Opposite directions (contradictory)

## 5. Use in Vector Databases

When data is stored as embeddings (numeric vectors), cosine similarity finds the most semantically similar items.

Steps:

1. Convert the query into a vector (embedding).
2. Compare it to all stored vectors using cosine similarity.
3. Return items with the highest similarity values.

### SQL Example (Postgres + pgvector):

```
SELECT id, embedding <#> query_embedding AS distance FROM items ORDER BY embedding <#>
query_embedding LIMIT 5;
```

Here, “`<#>`” is the cosine distance (1 - cosine similarity). Lower distance → higher similarity.

## 6. Advantages

- Ignores magnitude; focuses on direction.
- Works well for normalized embeddings.
- Ideal for semantic search in text, image, or recommendation systems.